

# Announcement

- Assignment 2 will be posted tonight
  - Due Jan. 23.

# Algorithm Performance

(the Big-O)

# Lecture 6

Today:

- Worst-case Behaviour
- Counting Operations
- Performance Considerations
- Time measurements
- Order Notation (the Big-O)

# Pessimistic Performance Measure

- Often consider the *worst-case* behaviour as a benchmark.
  - make guarantees about code performance under all circumstances
- Can predict performance by counting the number of “elementary” steps required by algorithm in the worst case
  - derive total steps ( $T$ ) as a function of input size ( $N$ )

# Analysis of dup\_chk ()

Q. What is  $N$ ?

- The number of elements in the array

```
int dup_chk(int a[], int length) {
```

```
    1  int i = length;
```

```
    N+1 while (i > 0) {
```

```
        N    i--;
```

```
        N    int j = i - 1;
```

```
        i+1    while (j >= 0) {
            i        if (a[i] == a[j]) {
                        return 1;
                    }

```

```
        i        j--;
```

```
    }
```

```
}
```

```
    1  return 0;
```

```
}
```

Outside of loop: 2 (steps)

Outer loop:  $3N + 1$

Inner loop:  $3i + 1$  for all possible  $i$  from 0 to  $N - 1$ .

$$= \frac{3}{2} N^2 - \frac{1}{2} N$$

Grand total =  $\frac{3}{2} N^2 + \frac{5}{2} N + 3$

*A quadratic function!*

$$1 + 4 + 7 + \dots + 3(N-1) + 1$$

# Some Math

$$1 + 4 + 7 + \dots + 3(N-1) + 1$$

# Some Math

$$1 + 4 + 7 + \dots + 3(N-1) + 1 = (3N-3+1 + 1) * N/2$$

# Some Math

$$\begin{aligned} 1 + 4 + 7 + \dots + 3(N-1) + 1 &= (3N-3+1 + 1) * N/2 \\ &= 1/2 * (3N-1) * N \end{aligned}$$



# Some Math

$$\begin{aligned} 1 + 4 + 7 + \dots + 3(N-1) + 1 &= (3N-3+1 + 1) * N/2 \\ &= 1/2 * (3N-1) * N \\ &= 1/2 * (3N^2 - N) \end{aligned}$$

# Some Math

$$\begin{aligned} 1 + 4 + 7 + \dots + 3(N-1) + 1 &= (3N-3+1 + 1) * N/2 \\ &= 1/2 * (3N-1) * N \\ &= 1/2 * (3N^2 - N) \\ &= 3/2 * N^2 - 1/2 * N \end{aligned}$$

# Some Math

$$\begin{aligned} 1 + 4 + 7 + \dots + 3(N-1) + 1 &= (3N-3+1 + 1) * N/2 \\ &= 1/2 * (3N-1) * N \\ &= 1/2 * (3N^2 - N) \\ &= 3/2 * N^2 - 1/2 * N \end{aligned}$$

# Some Math

$$\begin{aligned}1 + 4 + 7 + \dots + 3(N-1) + 1 &= (3N-3+1 + 1) * N/2 \\&= 1/2 * (3N-1) * N \\&= 1/2 * (3N^2 - N) \\&= 3/2 * N^2 - 1/2 * N\end{aligned}$$

Observation 1: The  $1/2 * N$  term doesn't matter very much

# Some Math

$$\begin{aligned} 1 + 4 + 7 + \dots + 3(N-1) + 1 &= (3N-3+1 + 1) * N/2 \\ &= 1/2 * (3N-1) * N \\ &= 1/2 * (3N^2 - N) \\ &= 3/2 * N^2 - 1/2 * N \end{aligned}$$

Arithmetic series



Observation 1: The  $1/2 * N$  term doesn't matter very much

Observation 2: Arithmetic series have  $N^2$  leading terms

1. As  $N$  gets bigger, the last number that we add is bigger
2. The number of pairs of numbers is bigger

# Analysis of dup\_chk ()

Q. What is  $N$ ?

- The number of elements in the array

```
int dup_chk(int a[], int length) {
```

```
    1  int i = length;
```

```
    N+1 while (i > 0) {
```

```
        N    i--;
```

```
        N    int j = i - 1;
```

```
    i+1    while (j >= 0) {
        i        if (a[i] == a[j]) {
                    return 1;
                }
    i        j--;
```

```
    }
```

```
    1  return 0;
```

```
}
```

Outside of loop: 2 (steps)

Outer loop:  $3N + 1$

Inner loop:  $3i + 1$  for all possible  $i$  from 0 to  $N - 1$ .

$$= \frac{3}{2} N^2 - \frac{1}{2} N$$

Grand total =  $\frac{3}{2} N^2 + \frac{5}{2} N + 3$


*A quadratic function!*

Observation: The  $\frac{5}{2} * N$  and 3 terms don't matter very much

# Empirical Measurement

- Another graph - a quadratic this time!
- Confirms predictions: **doubling** (2x) the input size leads to a **quadrupling** (4x) of the running time.

$N$	time (in ms)
10000	89
20000	365
40000	1424
100000	9011



# 2D Maximum Density Problem

Problem: Given a 2-dimensional array ( $N \times N$ ) of integers, find the  $10 \times 10$  swatch that yields the largest sum.



Applications:

- Resource management and optimization
- Finding brightest areas of photos





# Algorithm / Code?

- Simple approach: Try all possible positions for the upper left corner
  - $(N-9) \times (N-9)$  of them
  - use a nested loop
- Total each swatch using a 10x10 nested loop
- A *brute-force* approach!
  - Generate a possible solution [naively]
  - Test it [naively]

# In C

```
int max10by10(int a[N][N]) {  
    int best = 0;  
    for (int u_row = 0; u_row < N-9; u_row++) {  
        for (int u_col = 0; u_col < N-9; u_col++) {  
            int total = 0;  
            for (int row = u_row; row < u_row+10; row++) {  
                for (int col = u_col; col < u_col+10; col++) {  
                    total += a[row][col];  
                }  
            }  
            best = max(best, total);  
        }  
    }  
    return best;  
}
```

**x(N-9)**

**x(N-9)**

**x10**

**11**

**10**

**10**

Approximate Method:

Count the *barometer instructions*, the instructions executed most frequently. Usually, in the innermost loop.

Innermost loop:  $11 + 10 + 10 = 31$  ops

Total =  $31 \times 10 \times (N-9) \times (N-9) = 310N^2$

# Which Performance Measurement?

- Empirical timings
  - run your code on a real machine with various input sizes
  - plot a graph to determine the relationship
- Operation counting
  - assumes all elementary instructions are created equal
- Actual performance can depend on much more than just your algorithm!

# Running Time is Affected By . . .

- CPU speed
- Amount of main memory
- Specialized hardware (e.g., graphics card)
- Operating system
- System configuration (e.g., virtual memory)
- Programming Language
- Algorithm Implementation
- Other Programs
- . . .

# Comparing Algorithm Performance

- There can be many ways to solve a problem, i.e., different algorithms that produce the same result
  - E.g., There are numerous sorting algorithms.
- Compare algorithms by their behaviour for large input sizes, i.e., as  $N$  gets large
  - On today's hardware, **most** algorithms perform quickly for small  $N$
- Interested in growth rate as a function of  $N$ 
  - E.g., Sum an array: *linear* growth =  $O(N)$
  - E.g., Check for duplicates: *quadratic* growth =  $O(N^2)$

# Order Notation (the Big-O)

- Suppose we express the number of operations used in our algorithm as a function of  $N$ , the size of the problem.
- Intuitively, take the dominant term, remove the leading constant, and put  $O( . . . )$  around it.
- E.g.,  $f(N) = 348N^2 - 6956N + 34762 \rightarrow O(N^2)$

# Formalities of the Big-O

- Given a function  $T(N)$ , we say  $T(N) = O(f(N))$  if  $T(N)$  is at most a constant times  $f(N)$ , except perhaps for some small values of  $N$ .

- Properties:

- constant factors don't matter
- low-order terms don't matter

- Rules:

- For any  $k$  and any function  $f(N)$ ,  $k \cdot f(N) = O(f(N))$ 
  - E.g.,  $5N = O(N)$
  - E.g.,  $\log_a N = O(\log_b N)$ . Why?
  - Q. Do leading constants really not matter?

## Logarithm

From Wikipedia, the free encyclopedia

**Change of base** [edit] the

The logarithm  $\log_b x$  can be c

$$\log_b x = \frac{\log_k x}{\log_k b}.$$

# Leading Constants - Experiment

Of course, constant factors affect performance

- E.g., If two different algorithms run in  $f_1(N) = 20N^2$  and  $f_2(N) = 2N^2$  respectively, you expect Algorithm 2 to be 10 times faster.
- E.g., Similar algorithms can have the same running time.
- Big-O hides leading constants - a *hardware independent analysis*.



## Cray Supercomputer

17.6 x 10<sup>15</sup> instructions per second  
runs optimized dup\_chk( ) code from last time  
 $f(N) = 3/2 N^2 + 5/2 N + 3$

VS

## iMac Desktop Personal Computer (2011)

40 x 10<sup>9</sup> instructions per second  
runs an unoptimized, different dup\_chk ( )  
 $f(N) = 30N \log N + 5N + 4$



# Experimental Results

$N$	iMac	Cray
100,000	1.2 ms	850 ns
$10^6$	15 ms	85 $\mu$ s
$10^7$	0.2 s	8.5 ms
$10^8$	2 s	0.85 s
$10^9$	22 s	1.75 min
$10^{10}$	4.2 min	2:22 hr
$10^{11}$	56 min	10 days
$10^{12}$	8:20 hr	2.7 years



## Conclusions:

- Cray runs  $O(N^2)$  algorithm
- iMac runs  $O(N \log N)$  algorithm which runs faster than Cray for large  $N$  ( $10^9$  and beyond)
- Thus slow computer + no opt +  $O(N \log N)$  >> fast computer + optimization +  $O(N^2)$  algorithm
- **Rule of Thumb: The slower the function grows, the faster the algorithm.**
- For the  $O(N^2)$  Cray, a 10x increase in  $N$  leads to roughly a 100x increase in running time.
- For the  $O(N \log N)$  iMac, a 10x increase in  $N$  leads to roughly a 10x increase in running time (for the  $N$ ), plus a little (for the  $\log N$ ).

# Some Plots


$100\log(x)+60$  (yellow)

$0.3x+200$  (purple)

$0.18x\cdot\log(x)$  (green)

$0.0001x^2+0.2x+50$  (red)

$0.00000005x^3$  (black)



Courtesy of [fooplot.com](http://fooplot.com)

# Some Plots t


$100\log(x)+60$  (yellow)

$0.3x+200$  (purple)

$0.18x*\log(x)$  (green)

$0.0001x^2+0.2x+50$  (red)

$0.00000005x^3$  (black)



# What Does It All Mean?

- A carefully crafted algorithm can make the difference between a usable and an useless piece of software
- E.g., If it costs one algorithm 0.5s to search 1 billion bank records and another one 0.005s.
- E.g., Or, if  $10^9$  isn't "big" how about Google?
- E.g., Real-time systems - where a nearly instant response is required
- "You can't make a racehorse of a pig, but you can make a very fast pig."

# As N Gets Large, The Algorithm is Most Important

“You can’t make a racehorse out of a pig, but you can make a very fast pig.”

“When you find yourself trying all sorts of clever implementation tricks to speed up an algorithm, you should step back for a moment and ask if you’re trying to make a racehorse out of a pig. It might be more productive to use your time to look for a racehorse.”

“It’s important to know whether an efficient algorithm is possible. There’s no sense spending time looking for a unicorn. In this situation, the best you can hope for is a fast pig and a short racecourse.”

- Lou Hafer, SFU CS

# Big-O and Barometer Instructions

Rule of thumb:

The frequency of the barometer instructions will be proportional to the big-O running time

So, find the most frequent operation(s) and count them!

# Loops Multiply

```
int max10by10(int a[N][N]) {  
    int best = 0;  
    for (int u_row = 0; u_row < N-9; u_row++) {  
        for (int u_col = 0; u_col < N-9; u_col++) {  
            int total = 0;  
            for (int row = u_row; row < u_row+10; row++) {  
                for (int col = u_col; col < u_col+10; col++) {  
                    total += a[row][col];  
                }  
            }  
            best = max(best, total);  
        }  
    }  
    return best;  
}
```

**x(N-9)** (outermost loop)  
**x(N-0)** (middle loop)  
**x10** (innermost loop)  
**x10** (innermost loop)  
**barometer instructions** (points to the innermost loop body)

$$f(N) = 3 \times 10 \times 10 \times (N-9) \times (N-9) = O(N^2)$$

# Polynomials

Rule:

The powers of  $N$  are ordered according to their exponents, i.e.,  $N^a = O(N^b)$  if and only if  $a \leq b$

- E.g.,  $N^2 = O(N^3)$ , but  $N^3$  is not  $O(N^2)$ .

Why are lower-ordered terms not included?

- E.g., If your bank account followed  $f(N) = N^2 + N + 1$ , you would probably care a lot about the lower-ordered terms for small  $N$ , like  $N=5$ , as  $f(5) = 5^2 + 5 + 1 = \$31$ . You'll care about every dollar. But not for larger  $N$ , like  $N=1000$ , as  $f(1000) = 1000^2 + 1000 + 1 = \$1,001,001$ . You care most that you have that million bucks, and not much about the \$1000 or the \$1.



# More Rules

3. A logarithm grows more slowly than any other positive power of  $N$ .
  - E.g.,  $\log_2 N = O(N^{1/2})$ .
4. If  $f(N) = O(g(N))$  and  $g(N) = O(h(N))$  then  $f(N) = O(h(N))$ .
5. If both  $f(N)$  and  $g(N)$  are  $O(h(N))$  then  $f(N) + g(N) = O(h(N))$
6. If  $f_1(N) = O(g_1(N))$  and  $f_2(N) = O(g_2(N))$  then  $f_1(N) \times f_2(N) = O(g_1(N) \times g_2(N))$   
E.g.,  $(10 + 5N^2)(10\log_2 N + 1) + (5N + \log_2 N)(10N + 2N \log_2 N)$

# Typical Growth Rate Functions

- $O(1)$  – **constant** time
  - The time is independent of **N**, E.g., list look-up
- $O(\log N)$  – **logarithmic** time
  - Usually the log is to the base 2, E.g., binary search
- $O(N)$  – **linear** time, E.g., linear search
- $O(N \log N)$  – E.g., quicksort, mergesort
- $O(N^2)$  – **quadratic** time, e.g. selection sort
- $O(N^k)$  – **polynomial** (where  $k$  is a constant)
- $O(2^N)$  – **exponential** time, very slow!